

Introduction

COL331

Abhilash Jindal


Reference. OSTEP book: Chapter 2

Administrivia

- <http://abhilash-jindal.com/teach.html>
- Grading criteria, TAs, late policy, audit criteria, quizzes, labs, project, piazza link

Why does OS matter to a computer?

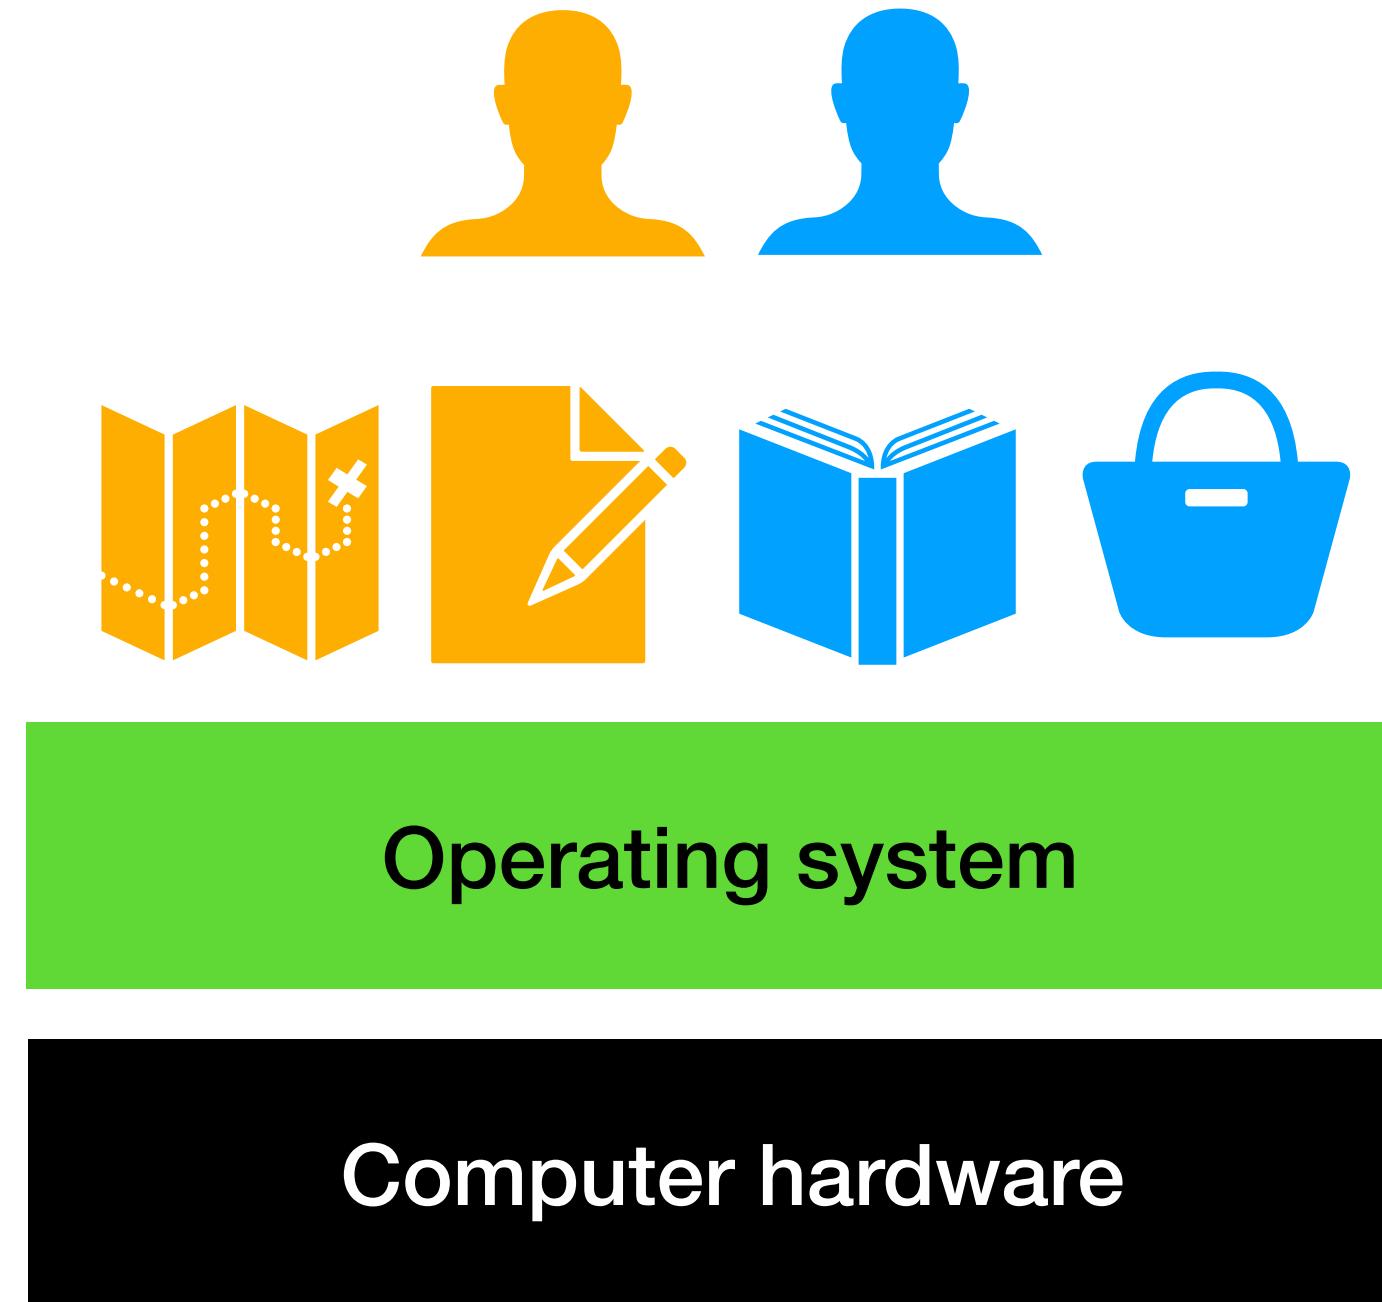
Computer organization



Purpose of an OS

- **Resource management**
- Provide higher-level services
- Protection and isolation

Purpose of OS: Resource management



- Example: `cpu.c`
- Give the illusion of more CPUs than there are
- Multiplex the hardware

Calculator analogy: Computing long sum



20
10
30
50
30
10
20
10


- $2 \ 0 =$ (move pointer to 10)
- $+ 1 \ 0 =$ (move pointer to 30)
- $+ 3 \ 0 =$ (move pointer to 50)
- $+ 5 \ 0 =$ (move pointer to 30)
- $+ 3 \ 0 =$ (move pointer to 10)
- $+ 1 \ 0 =$ (move pointer to 20)
- $+ 2 \ 0 =$ (move pointer to 10)

Sharing the calculator

20	10
10	70
30	20
50	40
30	20
10	10
20	50
10	10

- Steps to share the calculator:
 - $20 + 10 = 30 + 30 = 60$
 - Write 60 in notebook, remember that we were done till 30, give calculator
 - $10 + 70 = 80$
 - Write 80 in notebook, remember that we were done till 70, give the calculator back

Remember whatever is on the screen and give calculator?



- 2 0 = (move pointer to 10)
- + 1 0 = (move pointer to 30)
- + 3 0 = (move pointer to 50)
- + 5 0 = (move pointer to 30)
- + 3 0 = (move pointer to 10)
- + 1 0 = (move pointer to 20)
- + 2 0 = (move pointer to 10)

Can I give calculator here?

“Save 1 and remember pointer to be at 10”

No! Sum would be wrong!

“+ xx = (move pointer)” has to be atomic



Resource manager: multiplexing CPU

- CPU is also executing one instruction after another and incrementing “instruction pointer” *atomically*
- OS switches CPU between processes in the same manner as our calculator example (mechanism)
- What should happen when multiple processes want to run simultaneously? (policy)
 - Fairness: One banker got more calculator time than others
 - Often need to break away from fairness. Game should get more CPU time than Dropbox to provide good user experience
 - Starvation freedom: When there are multiple bankers, one banker never got the calculator

Purpose of OS: Resource manager

Different approaches to memory management:

- Segment different memory portions to different processes
- Multiplex memory pages across processes



Operating system

Computer hardware




Or



Memory management



- Segmentation is cheap to implement
- But not flexible. What if a process needs more memory than what OS gave?
- Paging is complicated to implement
- Highly flexible



Purpose of an OS


- Resource management
- **Provide higher-level services**
- Protection and isolation

Purpose of OS: Provide higher-level services



Operating system

Computer hardware




Example: io.c

Disk interface: List of blocks



File system OS interface: Folders and files

```
int fd = open("/tmp/file", O_WRONLY | O_CREAT);  
int rc = write(fd, "hello world\n", 13);  
close(fd);
```




Why file system?

Why not just multiplex disk blocks like memory?



Operating system


Computer hardware




- Disk blocks live after programs exits, computer restarts
- Different programs read / write same file
 - vim writes io.c
 - gcc reads io.c, write io
 - We finally run io

Higher level services provide portability

- Abstract away hardware details
- Programs need not be rewritten when moving from hard-disk drive to solid state drive



```
int fd = open("/tmp/file", O_WRONLY | O_CREAT);  
int rc = write(fd, "hello world\n", 13);  
close(fd);
```




Purpose of an OS

- Resource management
- Provide higher-level services
- **Protection and isolation**


Purpose of OS: Protection

- S_IRUSR | S_IWUSR: File can only be rw by user
- Disallow inappropriate accesses. Example: users reading each other's files




Purpose of OS: Protection

- Disallow inappropriate accesses.
Example: processes reading each other's memory



Purpose of OS: Isolation



Course structure: OS in action

- We will build an OS (xv6) from scratch
 - Booting: Bootloader, ELF format
 - Input-output: Programmable interrupt controllers, traps, interrupt descriptor table
 - File system: FS layout, buffer cache layer, name layer, crash consistency, devices as files
 - Processes: memory segmentation, rings, process table, context switching, scheduling, system calls, exec system call
 - Concurrency: data races, different types of locks
 - Memory virtualization: memory hierarchy, address translation mechanism, demand paging, thrashing, fork system call
 - Shell: Pipes, IO redirection
 - Parallelism: Enable more CPUs, revisit locks

Data races due to concurrency

./threads 100000


Thread 1	Thread 2
Read counter = 0	
Write counter = 1	
	Read counter = 1
	Writer counter = 2
Read counter = 2	
	Read counter = 2
	Writer counter = 3
Writer counter = 3	

**Why should I learn OS in 2026?
Isn't it a solved problem?**

We have indeed made good progress ..




Computer operators

1960s: Batch jobs



Personal computers

1980s: Interactive jobs!




UNIX.[®]

00011110 00011110 00011110 00011110 00011110 00011110




2000s: Cloud




- Cloud: I can rent *virtual machines* so I don't have to buy and manage servers.
- Hypervisor provides facilities to operating systems that OS provides to processes
 - multiplexes hardware among OS
 - protects and isolates OS from each other
- Hypervisors fundamentally enabled cloud computing

2000s: Distributed systems




- My data does not fit in a single machine. Distributed file systems; e.g., crawled web data is stored in Google File System
- My “tasks” span 100s-1000s of machines; e.g., computing page rank, training large language models over the entire internet



2010s: Smartphones

- New kinds of higher-level services: localisation, cellular, accelerometer, touch interface, etc.
- Resource constraints: Power management, UI system, etc.
- Even higher-level services: voice recognition, augmented reality, etc.
- Increased security concerns because of increase in sensitive data with rise of mobile banking, UPI etc. and because of moving devices






2010s: Cyber-physical systems

OS must not crash! Formally verified OS. Example: seL4


Typical progression of systems research

- Systems optimise for the “common case”
 - If common case changes, we need to rethink OS design
- Macro examples
 - Personal computers: batch jobs to interactive jobs
 - Smartphones: resource constraints, new sensors
 - Cyber physical systems: risk of human life




More trends: persistent memory

- Memory: volatile but fast
- Disk: persistent but large
- PMEM: persistent and fast!
 - PMEM aware file systems



More trends: big memory

- OS were designed when memory was scarce: few KBs
- You can now buy a server with 12TB of DRAM!
- Transparent huge pages
- Caching: Redis/memcached, etc.
- In memory compute: Spark, etc




More trends: fast networks

- OS typically assumed network is *much slower* than DRAM
 - Far memory

2020s	Latency	Bandwidth
DRAM	15ns	400 GBps
Ethernet	500ns	50 GBps

More trends: Rise of Unikernels




- OS optimises for common behaviours across all applications
- Each OS is now running only single application
- Unikernels optimise only for a single application

Confidential computing


- Developers want to run their proprietary code/models on sensitive data on someone else's machine (cloud) where all software including (host) OS is untrusted

Trusted Execution Environments (TEEs)




Hardware enclaves with Intel SGX

"I just trust my app code and the chip."




Hardware Confidential VMs with AMD SEV-SNP & Intel TDX

"Microsoft cannot touch my stuff in my VM."



Trusted launch VMs

"Only known, trusted code is running on my VM."



Trust

More trends: ML/LLM Systems

- ML training at scale
 - Pathways/TensorFlow (Google), Ray (Anyscale), ...
- LLM inference at scale
 - VLLM, SGLang, etc.
- Agentic programs




Figure 1: **Execution workflows for Agentic Programs.** Agentic programs are highly dynamic execution workflows that follow a directed acyclic graph (DAG). It consists of **LLM calls** from one or more LLM agents and **external interrupts** (i.e. tool calls, humans).

Image from Autellix paper

Why should I care about learning OS?

If I don't want to do systems research

- OS is a study of abstraction. Absorb all the complexity away from the developer / the user.
- OS has to provide high performance. Manage resources, provide isolation and protection with minimal overheads.
- Principles are useful when designing any practical system